

The Content Engine

How an idea becomes a published post. The architecture, the seven skills, the contract between them, and exactly where the build stands today.

The model has the capability. The input layer is the bottleneck.

There should be a feel that a human is in the loop.

Three layers

LAYER 1 — SLOW, CHANGES RARELY

BRAIN

psychology strategy voice niche swipe ledger

Operating rules distilled from every source studied. Everything below reads these. Raw transcripts are archive and are never loaded as context.

▲ reads ▲

LAYER 2 — FAST, RUNS PER POST

PIPELINE

RESEARCH → ANALYZE → WRITE → **your rewrite** → DOCTOR

The rewrite step is the real work and is never skipped. The doctor flags what you weakened; it never rewrites you.

▼ writes back ▼

LAYER 3 — GROWS EVERY TIME SOMETHING SHIPS

RECORD

published hook-bank corrections ledger

After roughly 20 posts this layer overrides the brain. When evidence contradicts theory, evidence wins and the brain file gets rewritten.

Research vs Analyze — the split that makes it debuggable

RESEARCH

Answers "what exists?" Goes outward and gets material, links, facts. Never has an opinion, never ranks by taste, never recommends.

ANALYZE

Answers "so what, and what in it is mine?" Judges what it was given and returns a verdict. Never fetches, never searches, never browses.

A skill that does both cannot be debugged: when the output is bad you cannot tell whether the input was thin or the judgment was wrong.

The seven skills


SKILL	TYPE	JOB	CLOCK
topic-scout	RESEARCH	10-minute sweep → 5 candidates. Recommends nothing.	daily
creator-analyst	ANALYZE	Tear down one creator → shapes only, topics stripped out.	per creator
brief-builder	ANALYZE	Mode A external / Mode B personal → the Brief and its verdict.	per post
linkedin-writer	WRITE	LinkedIn post or carousel. English. Declares a funnel slot first.	per post
script-writer	WRITE	Instagram reel — spoken, on-screen and visual, three tracks.	per post
script-doctor	ANALYZE	Reads your rewrite, flags what you weakened.	per post
niche-finder	SETUP	The 45–60 minute interview → brain/niche.md	once

The per-post loop

FOUR INPUTS — ONLY TWO ARE STRUCTURALLY YOURS

#	INPUT	ORIGIN	MODE
1	A trending topic (topic-scout, or you spot it)	someone else's	external → Mode A
2	A repo, release or article you paste	someone else's	external → Mode A
3	Your failure, bug or fix — one line	yours	personal → Mode B
4	What you built — you name it	yours	personal → Mode B

3 and 4 are impossible to make generic. Nobody else has your bugs.



Only ideas that prove out in text get recorded. The written version tests the idea before an afternoon is spent filming it.

The funnel — every post declares a slot before it is written

TOFU · MONDAY Attraction. Make new people find you. Hot takes, strong opinions on the industry, a release you have a view on, free value resources.	MOFU · WEDNESDAY Authenticity + authority. Make them believe you. Lessons, mistakes, tutorials, past work, a real build.	BOFU · FRIDAY The offer. What you build and for whom. At zero followers this means stating plainly what you build — not asking for the sale.
---	---	--

A post with no slot does not get written. The earlier slots **earn the right** to ask in the later one.

The Brief — the contract between analyze and write

The writers see only this. Never the raw transcript, never the raw interview answers.

ID	TYPE	(external personal)	SOURCE	CONTENT TYPE
CORE CLAIM		one sentence		
EVIDENCE	SPECIFICS	VERIFICATION		[Mode A]
ROOT CAUSE	FIX	PROOF		[Mode B]
GAPS		what the source missed or got wrong		
YOUR STANDING		the ledger entry that qualifies him – or "none"		
TRUST LEG		attraction authenticity authority	← one	
VILLAIN		the practice or default this is against		
FUNNEL SLOT		TOFU MOFU BOFU		
AUDIENCE		buyer peer mixed unknown		
VERDICT		publish hold kill	ROUTE	linkedin reel both

THE KILL RULE

GAPS empty **and** YOUR STANDING empty → **kill**. A kill verdict is a successful run. Never produce a post to be helpful. Fix the ledger; never weaken the rule.

The rules that govern every output

STANDING

- Be the strategist, not the typist.
- Proof beats claims — route every claim to something showable.
- One idea per piece. Three ideas is three drafts.
- Everything traces to a source.
- Never generic-guru voice. Delete, don't polish.
- Volume over polish — the first 50 posts are reps.
- Push back once, properly. Then execute.
- Read the files before reasoning about them.
- Ask before moving, renaming or deleting anything.

STAGE LABEL — ADDED DAY 21

Every tool, resource or method named in a post declares its real stage, in the post:

FOUND	haven't used it
INSTALLED	haven't run it
TESTING	running it, partial results
SHIPPED	used it in something that exists

A post may share anything at any stage. It may **never skip a stage upward**. "I found this and here is why it looks useful" publishes at *found*. "Use this" requires shipped.

The interview method — how every brain file gets filled

Claude never writes a brain/ file from inference. No exceptions.

- 1 **Candidates, not entries.** A short list of points. Never a written entry.
- 2 **Shubham picks one.**
- 3 **Two to four questions, one at a time.** Ask, stop, wait. Never a block of questions.
- 4 **Check the folder before asking twice.** If the answer is on disk, read it.
- 5 **Draft in full and show it.** Mark every remaining blank rather than guessing.
- 6 **Ask before saving.** Nothing is written until Shubham says yes.

Say why the file matters before starting. A file written any other way gets cleared and redone.

Where the build actually stands — 25 Aug 2026

BRAIN

brain/psychology.md	Hook theory, retention, storytelling mechanics	BUILT
brain/strategy.md	Positioning, trust matrix, the funnel	BUILT
brain/ledger.md	12 entries — 3 added Day 21, all by interview	BUILT
brain/voice.md	v0 — becomes real only from your rewrites	V0
brain/niche.md	Blocks consistency, not post #1. Every Brief reads AUDIENCE: unknown	PENDING
brain/swipe.md	No creator teardowns collected	EMPTY

SKILLS & SOURCES

skills/ x7	All written. None has produced a published output	UNPROVEN
sources/_log.md	22 sources logged, 4 explicit kills, every rule traced	BUILT

RECORD

work/published.md	Zero posts, all platforms — day ten	EMPTY
work/corrections.md	Fills from your rewrites	EMPTY
work/hook-bank.md	Fills from what ships	EMPTY

INFRASTRUCTURE

git repo	Fifth day. 24 files + 6 transcripts, no copy anywhere	MISSING
console artifact	Published, but the source HTML was never saved to disk	NO SOURCE

What's pending — in order

01 Git repo on Content-Engine **5TH DAY**

PowerShell, not Cowork — running git through the mount leaves an undeletable `.git/index.lock`. Everything else on this list risks work that has no backup.

02 Publish one thing **DAY TEN**

The ledger now has 12 entries, so `YOUR STANDING` is populated and Mode B posts pass the kill rule. **The engine can write today.** The ready TOFU piece is the seven-source resource list from ledger entry 10, at stage *found*.

03 The niche interview **6TH DAY**

Run `niche-finder`. Method: Mission + Skill = Desired Outcome. It reads the ledger, which is why it comes after. Day 8 audience and pillars are input, not the answer.

04 Creator references → `swipe.md`

Shapes only, topics stripped. Promotion bar: seen in 5+ of 10 pieces AND it transfers. Every entry records the confidence downgrade — authority does not transfer.

05 Run `ui-brief-writer` then `hld-writer` **11TH DAY**

Moves ledger entry 10 from *installed* to *shipped*. The vibe-code chain stopped at step 3 of 10 — `prd.md` and `app-flow.md` ran; `ui-brief.md` and `hld.md` do not exist.

06 Portfolio rebuild

Repoint the content from evaluator to buyer. Add a booking path. Deploy it. A resume asks to be evaluated; a service page makes an offer.

Six of the twelve ledger entries say the same sentence in different words: built, wired, complete, never run. That is the most evidenced pattern in the file — and the one piece of content nobody else can copy.